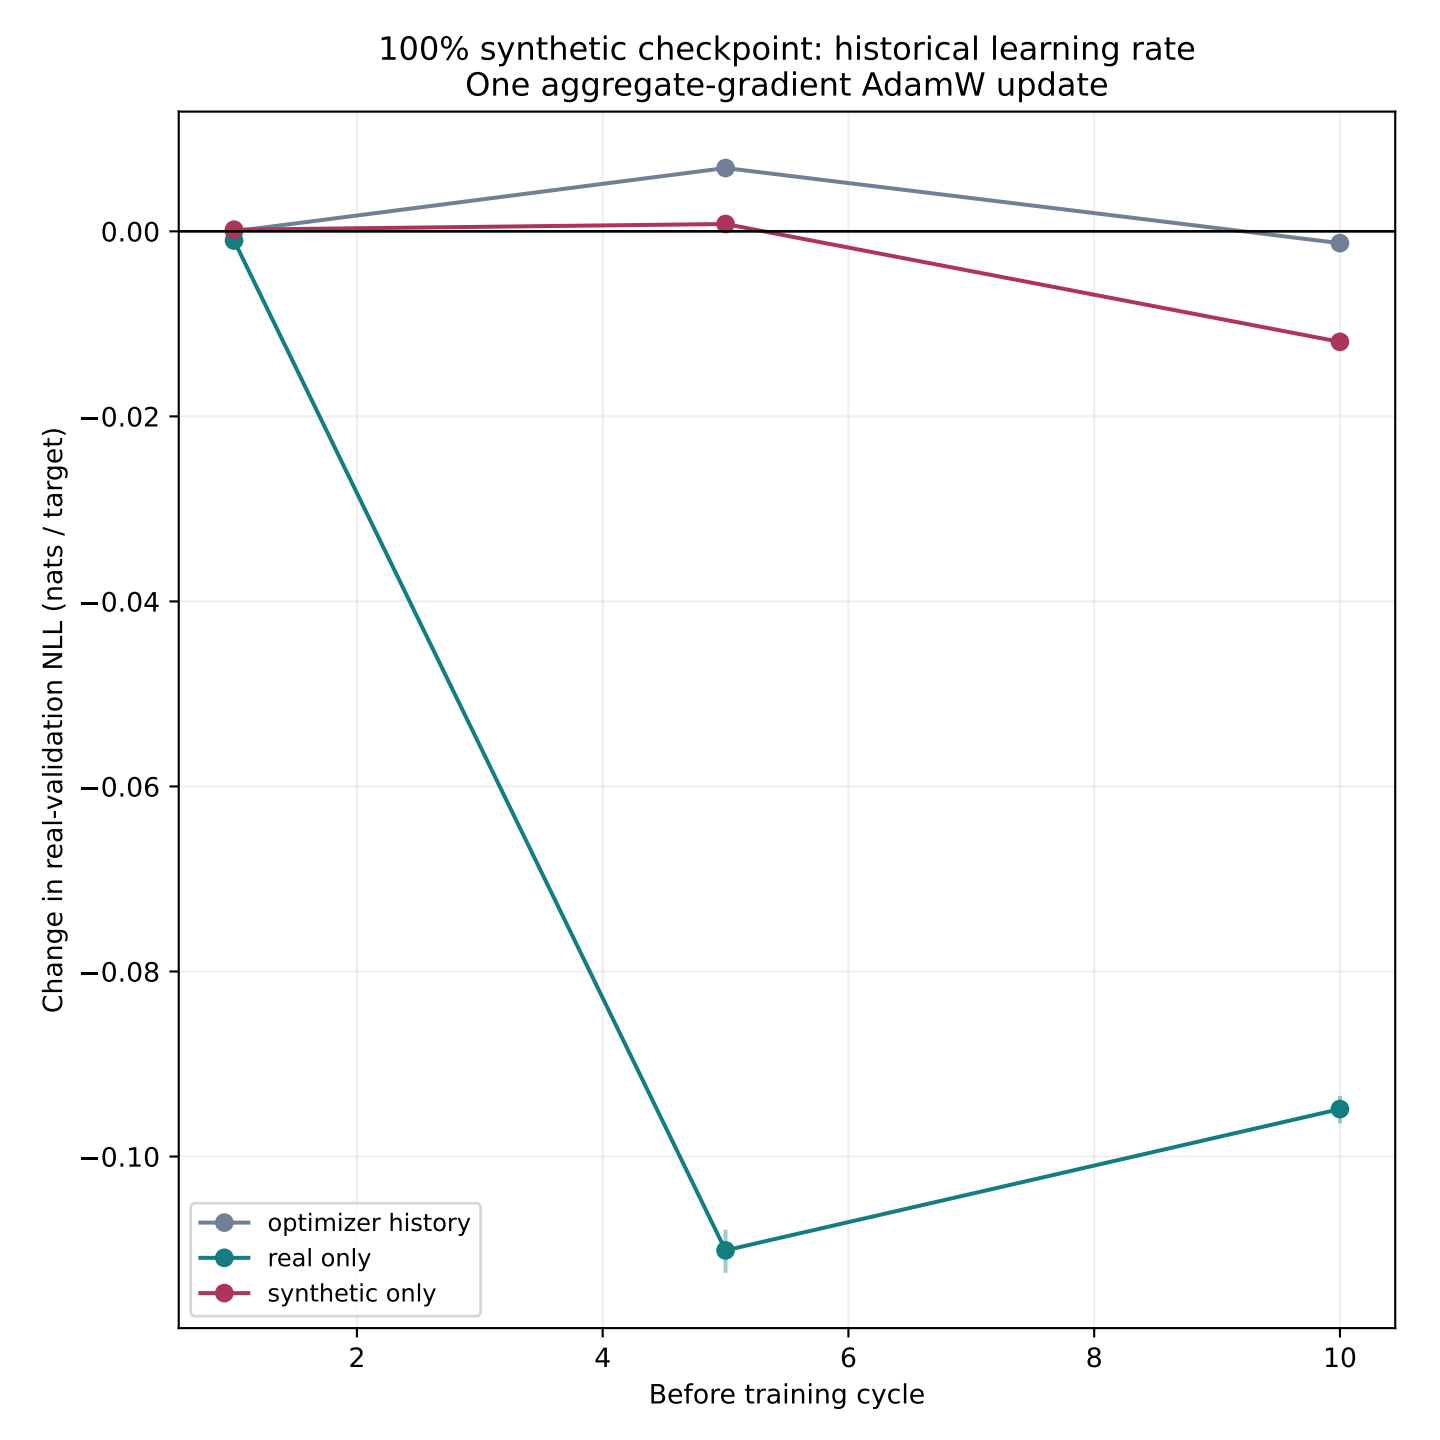
100% synthetic checkpoint: historical learning rate
One aggregate-gradient AdamW update
0.00
−0.02
−0.04
−0.06
−0.08
−0.10
Change in real-validation NLL (nats / target)
2
4
6
8
10
Before training cycle
optimizer history
real only
synthetic only