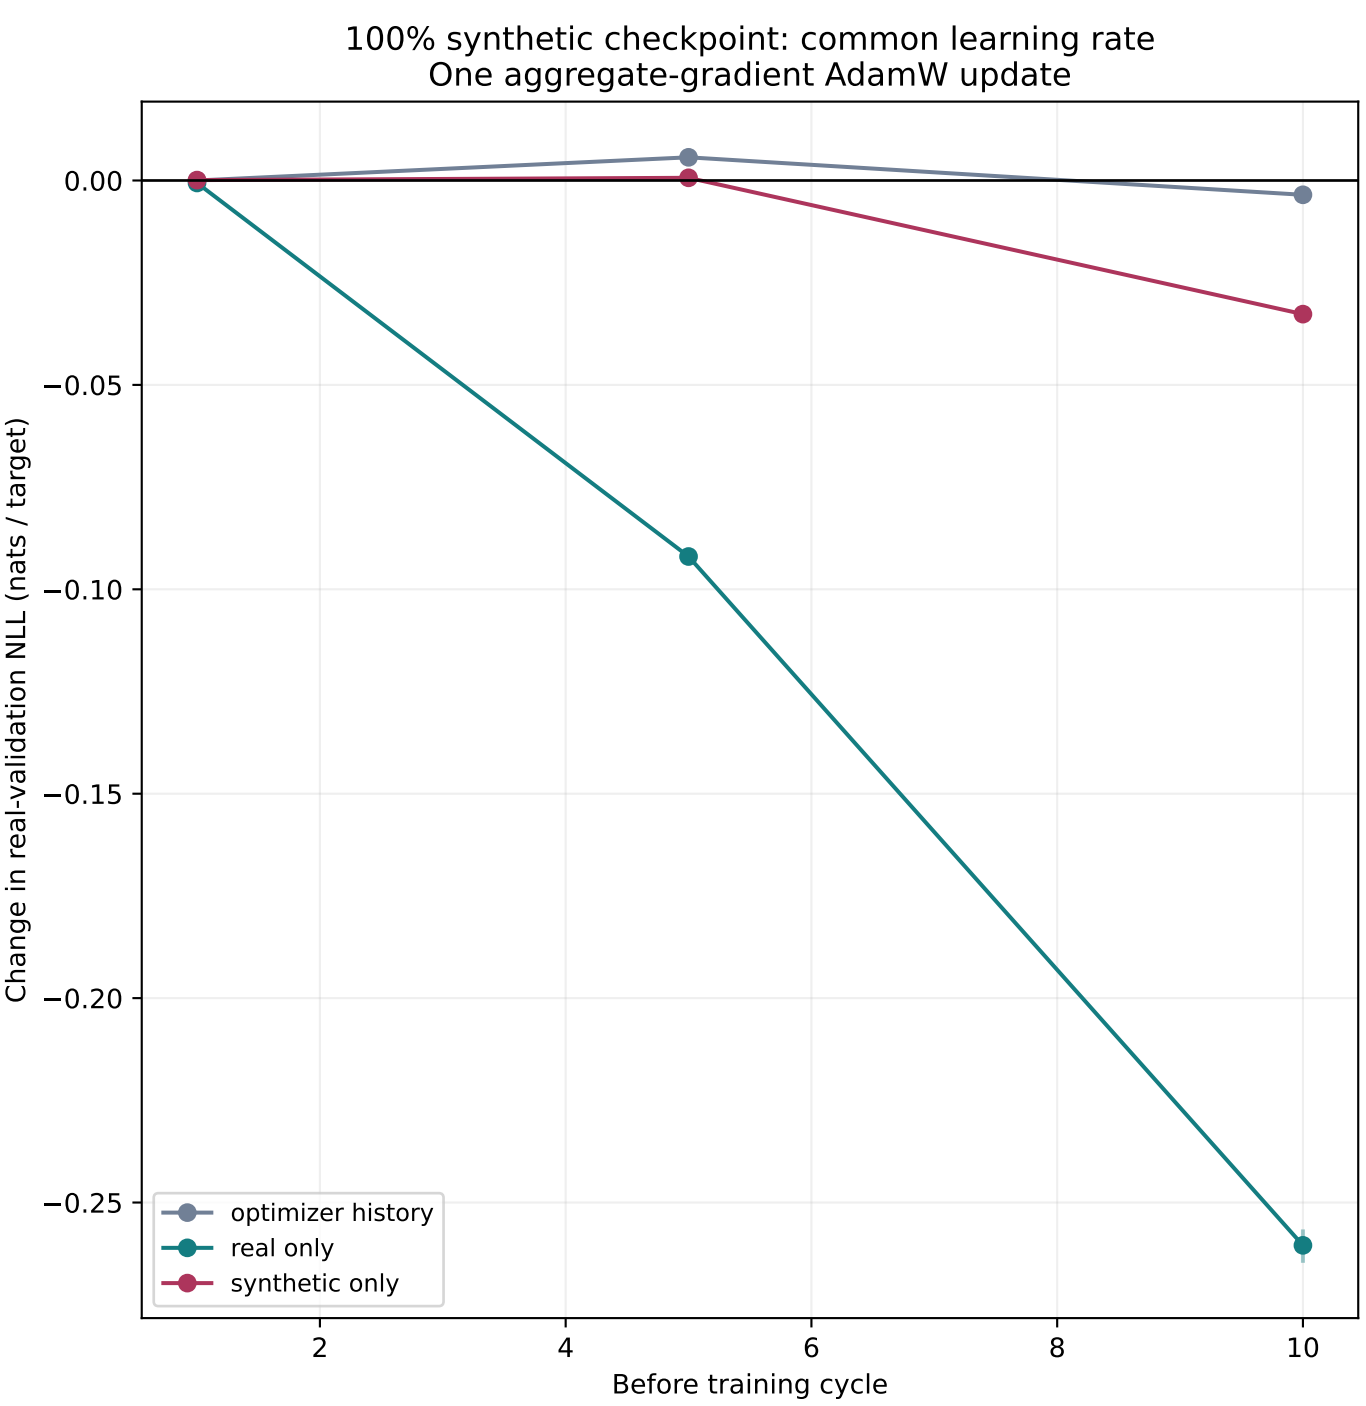
100% synthetic checkpoint: common learning rate
One aggregate-gradient AdamW update
Change in real-validation NLL (nats / target)
0.00
−0.05
−0.10
−0.15
−0.20
−0.25
2
4
6
8
10
Before training cycle
optimizer history
real only
synthetic only